

Killer Apps and Killer Robots

An Ethical Framework for Answering Questions of Automation

Peter McKay, Nisansa de Silva, Elizabeth Fuller, Bradley Green, Gautam Sondur
 Information and Computer science Department
 University of Oregon
 CIS 590 - Winter 2016

I. INTRODUCTION

"Sweating as he turned here and there to his bellows busily, since he was working on twenty tripods which were to stand against the wall of his strong-founded dwelling. And he had set golden wheels underneath the base of each one so that of their own motion they could wheel into the immortal gathering, and return to his house: a wonder to look at. These were so far finished, but the elaborate ear handles were not yet on. He was forging these, and beating the chains out" writes Homer in Iliad [1]. In fact, Iliad is not alone in ancient literature to describe such contraptions; mentions of automatons can be seen throughout other Greek texts [2], ancient Liezi text of China talks about an android singer presented to the fifth king of the Chinese Zhou Dynasty, King Mu of Zhou [3], the 11th/12th century Indian text Lokapannatti talks about robot guards and soldiers [4]. Even though one might question whether these are about actual incidents or not, these records make one thing certain. That is the fact that the idea of self-working machines is not a new concept to humanity. It is just an extension of the unquenchable human desire to make their lives easier; the desire that gave birth to the rudimentary tools that propelled us up the evolutionary tree, branching up and away from our animal brethren.

With the digital revolution spearheaded by the discovery of the transistor, we are now at a time when the very same stories, which one might brush off as fantasies when read in the above ancient texts, are slowly but surely becoming commonplace. In place of the auto moving tripods of Hephaistos in Iliad, we now have auto driving cars cre-

ated by Google [5] and Autopilot mode enabled cars by Tesla Motors [6]. Instead of stopping at the creation of singing automatons such as the one from the Liezi text, we have created automated systems that can compose music [7] and create artworks [8]. We are using drones and other unmanned vehicles to wage war. Fully automated security robots are emerging [9]. So it is rational to say that we are going far and beyond the wildest imaginations of our forefathers.

The discussion about artificial intelligence have grown to the level of debating whether a future artificial intelligence that is comparable to a human in all wakes of life (so that unless specified a human will not be able to distinguish between the said AI and another human [10]), is sapient or not. If we do agree that the AIs are sapient, from that agreement raises the question, "Is humanity using slave labor when it employees sapient robots to do tasks for them?" If a robot commits an offence that is punishable by law, who is liable? Is it the creators of the AI program or is a sapient AI able to take legal responsibility of its own actions? If it is the latter, what kind of rights would the AI have in a legal setting? If it is the former how is it different from the "Sins of the son" fallacy where parents are held liable for crimes committed by their children.

Like every scientific breakthrough before it, automation has created a dichotomy; one group consisting of optimistic scientists such as sir Arthur C. Clarke who claim "The only way of discovering the limits of the possible is to venture a little way past them into the impossible" [11], and the other consisting of scientists such as Stephen Hawking, Elon Musk, and Bill Gates who are rather pessimistic

about the advent of artificial intelligence and thus warns that humanity will have to be diligent as the technology develops [12]. Those who are not optimistic about the emergence of the automation also claim that it need not be developed to yield killer robots or artificial general intelligence [13] for automation to have a deep social impact. They clame that even a small amount of automation might have far reaching consequences. The objective of this paper is to analyze that social impact in an ethical standpoint and propose an ethical framework that can be used to mitigate the possible undesirable consequences of automation.

In section II we will discuss the socio-ethical impact that the industrial revolution had on the society given that fact that it is the only adequately documented closest approximate technological automation based transformative period of human history that we can compare this new contemporary transformation with smarter and almost autonomous machines. Then we discuss the situation that the society and technology currently stands in section III. Moving forward, in section IV we analyze some of the main solutions that various individuals and groups have proposed to handle the situation. Here we talk about the following solutions; The Hope of the Singularity, Bill Joy, Jason Lanier, The Amish, and Deep Green Resistance. For each of these solutions we discuss the pros and then we discuss the shortcomings of these solutions which makes solitary implementation of these existing methods inadequate and sometimes even unethical.

In section V we will discuss the solution that we propose in this paper. The solution is explained in terms of; Identifying Stakeholders, General Classes of Stakeholders, and Defining a Coherent Value System. We explain how our proposed framework can be used by way of the auto-driving car use case in section VI. Finally, the section VII gives the conclusions of the paper.

II. BACKGROUND

Later on in our paper, we will put forth an ethical framework for use in answering questions pertaining to automation and related technologies. As an

important part of that framework, we aim to analyze who may be affected by changes in the state of the art of technology, and how those effects may ripple out to change society at large. While these analyses would be much simpler were it not the case, the fact of the matter is that technological innovation is part and parcel of human society. Advances in technological innovation cannot be extricated from the social conditions which gave rise to such innovation, nor can they be separated from the social consequences which result from the adoption of that technology. If we wish to base our analyses in historical fact, it seems prudent to first review some basic history with an eye towards the factors we intend to consider. If we wish to see a preview of the impacts that automation will have on the workforce, and consequently the society at large, it seems prudent to look at the ways automation shaped society as its effects on the economy were first made manifest.

Let us consider, for a moment, the steam hammer. The steam hammer, a powered hammer for forging and stamping of metal parts, was invented in the mid-1800s as increasingly large and complex machines necessitated the fabrication of larger and heavier components. It made routine the construction of many works that, only years before, would have been considered marvels of engineering and concerted human effort. It contributed to the construction of great ships, the production of delicate clockwork, and a massive increase in the industrial productivity of Western Europe that saw an increase in the availability of cheaply manufactured goods formerly restricted to only the richest members of society. It also contributed [14] to a sharp decrease in the leverage that skilled engineers held over their employers, resulting in drastic consequences for working environments, and sharp increases in child-labor. James Nasmyth, one of those credited with the invention of the steam hammer, made the following statement regarding the power of automation in the modern (as of 1851) factory:

The characteristic feature of our modern mechanical improvements is the introduction of self-acting machinery. What every mechanical workman has now to do, and what every boy can do, is not to work himself, but to superin-

tend the beautiful labour of the machine. The whole class of workmen that depend exclusively upon their skill is now done away with. Formerly, I employed four boys to every machine. Thanks to these new mechanical combinations, I have reduced the number of grown-up men from 1,500 to 759. The result was a considerable increase in my profits. [15]

From the specific, we proceed to the general. The example of the steam hammer (in particular, the contrast between its positive and negative social effects) raises one of the first difficulties we will be forced to surmount in the course of developing our ethical framework. In the case of the industrialization of the workforce of Western Europe, those who wished to deploy greater degrees of automation in their factories were directly at odds with those who work in those factories. This is reflective of the effects of automation, not just on the market for the product they create, but also on the labor market. Labor, like any other commodity, has a price, which fluctuates based on supply and demand. When the owners of the factory can produce just as much value from less work, the demand for that labor goes down, and its price drops. Those workers have less leverage, and, as observed by Gaskell, their working conditions suffer.

When balancing the gains in technology and standards of living that have resulted from this march of progress, we see that there are two sides of this coin we must consider. The society as a whole may benefit, but what of the workers that toil in the factory? Does the gain outweigh the harm? Fortunately, in this particular case, we are not forced to make such a decision: social and cultural pressures intervened. The labor movement, in an internationally coordinated effort, slowly but surely pushed many governments into enacting legislation that protected workers from unfair practices on the part of their employers.

And yet, we are faced with a somewhat tougher questions, when we step back and consider the broader context around the industrialization of Western Europe. Industrialization is a somewhat titanic effort, it requires the investment of massive amount of capital to begin to see the advantages resulting from leveraging the economies of scale that make automation profitable. Different societies have achieved such a concentration of wealth

in different ways. China, in the course of The Great Leap Forward, began industrialization on a massive scale by nationalizing all production, taking central command of labor and distributing it according to a grand plan for industrialization of the workforce as a whole. The Soviet Union took a similar approach, nationalizing industry and placing it under central command, but then decentralized the industry geographically, exporting raw materials to nearby states as it absorbed them, thus creating a demand for the growth of industry in those formerly independent states that drove their growth further outward until they started bumping up against NATO countries. Western European countries accumulated wealth in privately held trading companies by establishing an early foothold in mercantile capitalism through the Triangle Trade. That is, they establishing trade routes between Europe, Africa, and the Americas, alternating in the shipment of slaves, raw materials (such as sugar cane grown by slaves), and manufactured goods (such as rum produced from last year's sugar cane).

In each of these situations, there exists a contradiction between the forces acquiring enough wealth to assemble an industrial base, and the people that had to give up that wealth. In China and the Soviet Union, the costs were paid by the citizens of those countries, and they reaped the benefits slowly as their standards of living increased fairly evenly. In Europe, the standards of living and availability of luxury goods skyrocketed much more rapidly: after all, the cost of automaton was offloaded to somewhere else: Africa and the American colonies. In the Soviet Union and China, the analysis looks simple: did the country gain more than it lost? China suffered a number of famines as the Great Leap Forward encountered organizational difficulties, but over a period of 30 years, the average lifespan of a Chinese citizen nearly doubled.

When, however, we look at a situation like British industrialization, we now find that we must take a look at a much more diverse selection of groups in analyzing the impact of technology on the society. As globalization has blurred the lines between nations, so too has information technology

led to a vastly more interconnected world than we have ever seen. A change in technology today can lead to nearly instantaneous changes through enormous swaths of society. In the immediately following sections, we will further examine the current state of the world as it has evolved, the approaches that some groups take in analyzing that world, and the steps they deem necessary due to such an analysis.

III. THE CURRENT SCENARIO

On the 12th of March 2016, history was made. Lee Sedol, a professional Go Player who some consider to be the best player of the last decade, lost a game of Go. His third consecutive loss in a best-of-five competition. His opponent? AlphaGo an Artificial Intelligence based Go playing computer system. Although this wouldn't be considered anything new, with computers outsmarting their human opponents from Watson, IBM's Artificial Intelligence besting previous champions to win Jeopardy to Deep Blue (another of IBM's attempts at Artificial Intelligence), the first computer to beat a Grand Master at Chess. But why was this such a historic victory? It was for two primary reasons; firstly the game of Go has about the same number of possible positions as there are atoms in the universe (this rules out simply computing every possible combination) and secondly, this is a milestone in the trend of using computers to replicate mental processes that were earlier considered to be the sole dominion of the living beings, particularly humans.// As mentioned earlier, we live in the Information Age, and changes have a tendency to propagate through vast sections of society fairly quickly. Globalization and the subsequent transfers of technology enabled by the Internet, allows for a proliferation of ideas and techniques that would have been unheard of a few decades ago. Automation's strongest argument has been the one to reduce manual labor and improve the quality of human lives. Although now, it would seem that this is no longer the case. Many proponents of automation have pointed to the migration of people from unskilled and semi-skilled labor to more skilled labor which rely on mental processes that

could not have been automated at the time. It is evident that this sacred boundary is under threat of being or has already been, depending on who you ask, infringed upon.// The advent of Applications that have improved accessibility to services, from planning your week to simply setting an alarm or a reminder to do your taxes, all these tasks have been automated already. The automation of these activities begs the question whether this would make humans obsolete in performing these activities. The public sentiment does not seem to reflect this. In a study carried out by the *Pew Research Center* [16], although about two-thirds of Americans believe that computers and robots will be performing most of the work being done by humans in the next 50 years, nearly 80 % believe that their current jobs will still be around then.

A likely answer to this is skewed figure is probably to do with the General outlook most people have on technology. In a 1996 *Washington Post/Kaiser Family Foundation/Harvard* poll, 70% of Americans said the increased use of technology in the workplace was good for the economy. In a 2010 *Allstate/National Journal* poll, 79% said that information technology was extremely or very important to creating economic growth in the U.S. Some would say they aren't wrong but it should be noted that it is rare that general public opinion is aligned with the most ethical option.

Even though AI forms only a small section of this problem, it engenders the primary difference between the *First Industrial Revolution* and now. Technology Heavyweights like Elon Musk and Stephen Hawking have openly spoken about the perils of the technology. With some even saying it would lead to the extinction of Humans altogether because we would simply not live up to the standard of the subsequent 'higher- intelligence'. The extreme views on this issue are well documented, and we will further elaborate on them in the next section, but our aim remains to balance these views to ensure a smooth transition. There is no escaping the automation, the best we can do is ensure it is done in the most ethical fashion, without favoring one stakeholder over the other and ensuring that we have a holistic view of the situation rather

than proposing solutions that would be suited to a vacuum.

IV. A REVIEW OF PRESENT SOLUTIONS

A. *The Hope of the Singularity*

B. *Bill Joy*

C. *Jason Lanier*

D. *The Amish*

E. *Deep Green Resistance*

Deep Green Resistance is a group primarily identified by their beliefs and behavior with regard to specifically environmental matters, rather than more general opinions about technology. Specifically, DGR subscribes to the notion that industrial civilization poses too great of a threat to life on the planet earth, both human and otherwise, to be allowed to survive. The ultimate goal of DGR is the destruction of industrial civilization and a return to an earlier stage of societal and economic development.

DGR's philosophy draws from the Deep Ecology movement, which holds that an anthropocentric analysis of ecology, defining other forms of life in terms of their utility to humankind, does not do an adequate job of describing the complexity of ecological systems. DGR takes this a step further, in declaring all life to be equal to human life. Starting at the problem of inequality among life, they turn to ideals similar to the political stance of anarcho-primitivism: the idea that human society, industrial capitalism in particular, cannot be reformed into a more beneficial form.

Such a philosophy stems from an agreement with Friedrech Engal's seminal anthropological work on early civilizations [17], specifically the assertion that early gatherer-hunter societies were unable to produce structural inequality without the ability to acquire surplus wealth. Unlike Engals, modern Marxists, and even their fellow anarchists, DGR and anarcho-primitivists do not believe that a society can ever reach a level of egalitarianism and freedom from oppression that they are comfortable with. Rather, they believe that the only truly equal societies that have ever existed were

only able to function without the trappings of civilization. Therefore, they are willing to advocate any means to right the wrongs they see with the world. To them, technology is not evil per se. Rather, technology is another product of a system with which they cannot make peace, and as such, it must fall by the wayside.

Later in our paper, we advance an ethical framework for the analysis of decisions regarding technology and automation. In so doing, we emphasize that we must look to the past when examining possible consequences of an action, basing, wherever we can, our reasoning on historical fact. In this case, we would be remiss if we did not qualify our description of the views of DGR with the caveat that their views ultimately reduce to something chillingly familiar. A return to pre-industrial levels of technology would necessitate either the death of billions, or the restriction of new births to a fraction of their current levels. Without modern medical technology, it becomes a disturbing fact that disabled people will not be able to survive in such a world: this means DGR endorses either genocide or eugenics.

V. A COHERENT ETHICAL FRAMEWORK

A. *Identifying Stakeholders*

Ethical Analysis of an action, circumstance or rule must always begin with Stakeholder Analysis. Stakeholders are those who are affected as consequence of the action, circumstance or rule. In other words, the parties concerned have a stake in the eventual outcome of performing the action or implementing the rule.

The framework aims to take into account two main factors; firstly, factors that affect the identification of stakeholders; and secondly, when viewing automation we observe three general classes of stakeholders, with varying levels of interplay seen between these classes based on the aforementioned factors.

Factors affecting Stakeholder Identification:

The process of identifying Stakeholders is generally specific to the action and rule as well as to the circumstances that surround the process. Although

our framework strives to espouse the Kantian Ideal of universality, it would be naive to assume that the stakeholders would be a constant across the board; taking into account that our framework aims to perform an ethical analysis on the effect of automation. To this end, in order to identify the stakeholders we must try to obtain a holistic view of the situation. In order to do this, we enumerate three factors, which would have an effect in the identification process:

- **Cultural Bias:**

Cultural Bias has long hindered the adoption of technology. Rather than attributing this as a drawback (which it is sometimes) see it as nothing but the decision making process of a collective, based on previous experience as well as skepticism about new and sometimes untested methods/technologies. Like a double-edged sword, the same cultural bias might catalyze the adoption of the technology as well. This can be seen in the early problems faced during adoption of disruptor technologies like the Radio and aircrafts. In order to correctly identify stakeholders, it is imperative to accurately weigh the cultural bias associated with a given society; especially since the effects of the technology might be experienced only by a certain section of society.

- **Value System:**

Value systems play a big part in helping the acceptance of technology. Say for example, in the United States the Apple vs. FBI debate has tested the value systems of the public, where protection of individual liberties is of paramount importance, if a similar case would have taken place in the People's Republic of China, the prevailing value systems would have ensured massive public support for the Government's stand. In short, the classification of stakeholders would change dramatically where in one society, the Government represents the interests of the public where as in another, the individual/public takes on an identity of their own.

- **Economic Status:**

The Economic Status of a society broadly refers to the distribution of wealth in the soci-

ety. We see vastly contrasting structures in Developed Countries and in Developing Countries. One example that could elucidate this difference is the adoption of automation of jobs such as customer service. This would have varying effects depending on the country. In developed countries, these are considered tertiary level jobs and do not form the backbone of the economy. Hence the effect would be largely mitigated and with government intervention could also be avoided. But in the case of developing countries, where semi-skilled jobs like freight services employ a large number of people, we would see a larger effect in terms of unemployment. Accounting for this would play a large part in correctly weighing each stakeholder.

B. General Classes of Stakeholders

Taking into account the effect of automation over the entire cross-section of society, we think it is prudent to identify certain classes of stakeholders which would be present across any situation, regardless of the geographical location, economic status and value system. Although there would be different interplays between these classes given the above mentioned factors that affect the selection of stakeholders, their existence is no less in doubt by our standard.

It is important to add that although we define these classes, they are in no way exhaustive and the due diligence required to identify stakeholders must be done in order to ensure the holistic view we would like to achieve in this part of the process.

- **Public:**

The Public, refers to society at large who no matter what the situation will be directly/indirectly affected by the implementation of the automation. Typical issues include job/labor concerns, concerns regarding misuse against other members of the public, and concerns regarding laws enforcing the ethical use of the same, just to name a few. By far this is one of the larger classes, and within it many more might emerge, especially when the technology is targeted to a certain section/strata

of society. It is therefore essential to identify which section might have a larger stake in the implementation of the automation, be it positive or as it is often seen, negative.

- **Government:**

Government, refers to the regulatory body that oversees the ethical use of the automation. We can say in certain scenarios, the lines discerning public interest and government interest might blur, we considered these to be exceptional cases where we can identify one class to represent the interests and stakes of another (as mentioned in the example above). It can be seen in most modern societies today, the public does not necessarily gain adequate representation of their interests through government and similarly, the Government must at times ensure the security and well-being of the state even if that does not represent what the public deems fit. For this and reasons that outline similar cases of dichotomy, we would like to represent this as a separate class.

- **Corporate Interests:**

Automation today, as it was earlier as well, is driven by commercial interests. We represent these interests through the class of corporate stakeholders. These would represent the individuals as well as organizations which have driven the case for the automation and would benefit in some monetary fashion by the said implementation. Apart from this, we would like classify industries in this class of stakeholders as well, because of the widespread transmission of ideas and technology, practices linked with an increase in productivity would be widely adopted even if that is at the cost of jobs. Although it would be easy to classify corporate interests as strictly monetary, but in certain scenarios we see that this stakeholder has aligned interests with another classes (such as Government) and hence would require a somewhat different approach to weighing their stakes.

C. *Defining a Coherent Value System*

D. *Synthesis*

VI. SELF-DRIVING CARS: AN APPLICATION OF OUR FRAMEWORK

A. *Background*

Long a fascination for inventors, the seed for autonomous cars was sewn in the 1960s with the Stanford Cart [18]. The cart navigated small spaces by taking photos of its surroundings and using a simple computer program to analyze the photos and choose a path of travel. Two decades later, a vision-guided Mercedes-Benz van designed by Ernst Dickmanns navigated empty streets at speeds of up to 39mph. By 2009, Google began privately testing anonymous vehicles on closed streets, and by 2012, changes in various states' laws allowed for testing on city streets. As of 2016, Google claims over 1.3 million autonomous miles driven, although Google's cars are limited to 25mph and must always have a human driver present. [19] As driverless car technology improves and becomes more widely implemented, the true societal impact of this technology will begin to be felt.

B. *Application*

We're in an interesting time regarding autonomous vehicle technology. It's clear that the technology will become widely used, but it has yet to actually happen. This presents an opportunity to consider how, as a society, we should approach the ethical ramifications of this potentially society-altering technology. To do so, we'll apply the ethical framework laid forth earlier in this paper to autonomous cars.

To begin, it is important to identify the three main stakeholders in driverless car technology: the public, the government, and corporations. These stakeholders are universal, but the beliefs of each stakeholder can vary depending on cultural biases, value systems, and economic status. Also, the weight given to each stakeholder can vary based on these factors. In the interest of concision, this application will focus on a generalized idea of American cultural bias, values, and economic status.

Because the United States is such a demographically, economically, and culturally diverse country, it is difficult to make generalizations of the country as a whole. That being said, it is possible to extrapolate mainstream value systems of our three stakeholders (public, government, and corporations). The American public values individualism and personal rights. Corporations value free market principles and the acquisition of capital. The government acts as a mediator (sometimes poorly) between the interests of corporations and of individuals. Each stakeholder has the power to act on the other and bring about change. With the stakeholders, value systems, and power structure identified, it is possible to apply our ethical framework to autonomous vehicle technology in the United States.

Autonomous vehicles have a lot to offer a utilitarian. A recent Virginia Tech study [20] shows that the crash rate of Google's self-driving cars is lower than the national crash rate of conventional cars. Google's cars experience 3.2 crashes per million miles while conventional cars experience 4.2 crashes per million miles. With autonomous vehicle technology continually improving, the technology will become even safer in the years to come. This sounds great, but is a full on, no-strings-attached embrace of the technology the best course of action? For a utilitarian, the answer may be "yes", but there is a problem with this approach. Autonomous vehicles will face a variety of moral dilemmas, and a one-size-fits all ethical framework doesn't cover all of these possibilities sufficiently.

Consider the following example: during the November 2015 terror attacks in Paris, ride-sharing app Uber's algorithmic design dramatically raised prices in the areas near the attacks due to an increase in ride requests from people trying to flee the chaos. In this situation, stakeholders (public and corporate) were in direct opposition to each other. Uber saw the situation as a moneymaking opportunity, while the public saw Uber as price gouging and profiting from a disaster. Eventually, Uber gave in to public demand and lowered prices. Situations such as this will become the norm in the age of autonomous cars. In fact, Uber is aggressively

seeking driverless car technology for use in their taxi company.

In the case of Uber raising the prices on people fleeing a disaster, utilitarianism supports profit maximization as a means of supporting the greatest societal good. Clearly, this is a morally unfavorable position for the public, but fortunately, the moral framework set forth in this paper allows an exception to utilitarianism on the basis of protecting an individual's welfare and well-being. In the case of the Paris attacks or hypothetical future situations involving Uber autonomous vehicles, our moral framework prioritizes lives in danger over profits in dangers.

Consider another hypothetical example: you're riding down the highway in your autonomous car and a large object falls off of a truck in front of you. Your car cannot stop itself in time and must decide on the best course of action. If it continues straight and runs into the object, it places your life in serious jeopardy. There are two motorcyclists on either side of your car; one is wearing a helmet and the other is not. Swerving into one of the motorcyclists will likely save your life, but jeopardize the life of one of the motorcyclists.

With strict utilitarianism, your car would likely swerve into the motorcyclist wearing the helmet because it creates the best chance of survival for all involved, but is it fair to punish a motorcyclist for being responsible and wearing a helmet? What are the legal and moral ramifications of such a programmed decision? In this scenario, our ethical framework would fall back onto the autonomy of the individual. They could decide how their car is programmed to react in cases such as these. Alternatively, self-driving cars could rely on an element of randomness in making such decisions or even cede control to a human driver (local or remote) despite the car's ability to handle situations such as this without human intervention. The important point is leaving autonomy in the hands of the individual, not the machine.

While far from foolproof, our ethical framework allows for important exceptions to utilitarianism. The implementation of driverless cars will undoubtedly lead to many unique moral dilemmas,

therefore the existence of *fallback rules* to a society's moral framework will be critical. In the previous examples, the *fallback rules* of our ethical framework placed health and well-being above financial profit and the autonomy of a human over the autonomy of a machine. These are only two examples of the multitudes of moral dilemmas that will become ever-present in a driverless car future, and it is critical that we continue to work toward developing an ethical framework to deal with every possibility.

VII. CONCLUSION

Rise of the automated tools designed to accomplish certain tasks and ethical problems stemming thereof is a very potent area for dialogue. It is evident that the set of current solutions have more cons than pros when one put them under scrutiny. We claim that the framework that we proposed in this paper would minimize the short comings and maximize the gains. But one must always remember that until we do obtain a certain level of automation, real the social impact of such a technology can only be speculated. Minute details that escape the most prudent analysis might have very far reaching consequences. And thus we stand on the precipice and stare in to the void, and does the void stare in to us back?

REFERENCES

- [1] Homer, *Iliad*, 1194 BC, Chapter 18. 371 ff.
- [2] Aaron J. Atsma, "Automatones," 2015.
- [3] Lie Yukou, *Liezi*, pp. 83–86, 600 BC.
- [4] Sarah L. Higley, "Alien intellect and the roboticization of the scientist," *Camera Obscura*, vol. 14, no. 1-2 40-41, pp. 129–160, 7 1997.
- [5] "Google self-driving car project," 2016.
- [6] "Your autopilot has arrived," 2016.
- [7] I. Liu and B. Ramakrishnan, "Bach in 2014: Music Composition with Recurrent Neural Network," *ArXiv e-prints*, Dec. 2014.
- [8] James Gould-Bourn, "Man combines random people's photos using neural networks and the results are amazing," 2016.
- [9] Justin Mahboubian-Jones, "Fully-automated security robots have arrived," 2016.
- [10] Eyal Reingold, "The Turing Test," 1999.
- [11] Arthur C. Clarke, *Profiles of the Future*, chapter "Hazards of Prophecy: The Failure of Imagination", Phoenix, 1962.
- [12] Michael Sainato, "Stephen hawking, elon musk, and bill gates warn about artificial intelligence," 2015.
- [13] Luke Muehlhauser, "What is agi?," 2013.
- [14] Peter Gaskell, *The manufacturing population of England, its moral, social and physical conditions*, Baldwin & Cradock, 1833.
- [15] Karl Marx, "The strife between workman and machine," in *Capital*, vol. 1, chapter 15. 1867.
- [16] Aaron Smith, "Public predictions for the future of workforce automation," 2016.
- [17] Friedrich Engels and Tristram Hunt, *The Origin of the Family, Private Property and the State*, Penguin UK, 2010.
- [18] Hans P Moravec, *The Stanford cart and the CMU rover*, Springer, 1990.
- [19] Alex Davies, "Google's self-driving car caused its first crash," *Wired*, 2016.
- [20] Myra Blanco, Jon Atwood, Sheldon Russell, Tammy Trimble, Julie McClafferty, and Miguel Perez, "Automated vehicle crash rate comparison using naturalistic data," Tech. Rep., Virginia Tech Transportation Institute, 2016.