Killer Apps and Killer Robots
An Ethical Framework for Answering Questions of Automation

Peter McKay, Nisansa de Silva, Elizabeth Fuller, Bradley Green, Gautam Sondur

Information and Computer science Department

University of Oregon

CIS 590 - Winter 2016

I. INTRODUCTION

"Sweating as he turned here and there to his bellows busily, since he was working on twenty tripods which were to stand against the wall of his strong-founded dwelling. And he had set golden wheels underneath the base of each one so that of their own motion they could wheel into the immortal gathering, and return to his house: a wonder to look at. These were so far finished, but the elaborate ear handles were not yet on. He was forging these, and beating the chains out" writes Homer in Iliad [1]. In fact, Iliad is not alone in ancient literature to describe such contraptions; mentions of automatons can be seen throughout other Greek texts [2], ancient Liezi text of China talks about an android singer presented to the fifth king of the Chinese Zhou Dynasty, King Mu of Zhou [3], the 11th/12th century Indian text Lokapannatti talks about robot guards and soldiers [4]. Even though one might question whether these are about actual incidents or not, these records make one thing certain. That is the fact that the idea of self-working machines is not a new concept to humanity. It is just an extension of the unquenchable human desire to make their lives easier; the desire that gave birth to the rudimentary tools that propelled us up the evolutionary tree, branching up and away from our animal brethren.

With the digital revolution spearheaded by the discovery of the transistor, we are now at a time when the very same stories, which one might brush off as fantasies when read in the above ancient texts, are slowly but surely becoming commonplace. In place of the auto moving tripods of Hephaistos in Iliad, we now have auto driving cars cre-ated by Google [5] and Autopilot mode enabled cars by Tesla Motors [6]. Instead of stopping at the creation of singing automatons such as the one from the Liezi text, we have created automated systems that can compose music [7] and create artworks [8]. We are using drones and other unmanned vehicles to wage war. Fully automated security robots are emerging [9]. So it is rational to say that we are going far and beyond the wildest imaginations of our forefathers.

The discussion about artificial intelligence have grown to the level of debating whether a future artificial intelligence that is comparable to a human in all walks of life (so that unless specified a human will not be able to distinguish between the said AI and another human [10]), is sapient or not. If we do agree that the AIs are sapient, from that agreement raises the question, "Is humanity using slave labor when it employees sapient robots to do tasks for them?" If a robot commits an offence that is punishable by law, who is liable? Is it the creators of the AI program or is a sapient AI able to take legal responsibility of its own actions? If it is the latter, what kind of rights would the AI have in a legal setting? If it is the former how is it different from the "Sins of the son" fallacy where parents are held liable for crimes committed by their children.

Like every scientific breakthrough before it, automation has created a dichotomy; one group consisting of optimistic scientists such as sir Arthur C. Clarke who claim "The only way of discovering the limits of the possible is to venture a little way past them into the impossible" [11], and the other consisting of scientists such as Stephen Hawking, Elon Musk, and Bill Gates who are rather pessimistic

about the advent of artificial intelligence and thus warns that humanity will have to be diligent as the technology develops [12]. Those who are not optimistic about the emergence of the automation also claim that it need not be developed to yield killer robots or artificial general intelligence [13] for automation to have a deep social impact. They clame that even a small amount of automation might have far reaching consequences. The objective of this paper is to analyze that social impact in an ethical standpoint and propose an ethical framework that can be used to mitigate the possible undesirable consequences of automation.

In section II we will discuss the socio-ethical impact that the industrial revolution had on the society given that fact that it is the only adequately documented closest approximate technological automation based transformative period of human history that we can compare this new contemporary transformation with smarter and almost autonomous machines. Then we discuss the situation that the society and technology currently stands in section III. Moving forward, in section IV we analyze some of the main solutions that various individuals and groups have proposed to handle the situation. Here we talk about the following solutions; The Hope of the Singularity, Bill Joy, Jason Lanier, The Amish, and Deep Green Resistance. For each of these solutions we discuss the pros and then we discuss the shortcomings of these solutions which makes solitary implementation of these existing methods inadequate and sometimes even unethical.

In section V we will discuss the solution that we propose in this paper. The solution is explained in terms of; Identifying Stakeholders, General Classes of Stakeholders, and Defining a Coherent Value System. We explain how out proposed framework can be used by way of the auto-driving car use case in section VI. Finally, the section VII gives the conclusions of the paper.

II. BACKGROUND

Later on in our paper, we will put forth an ethical framework for use in answering questions pertaining to automation and related technologies. As an important part of that framework, we aim to analyze who may be affected by changes in the state of the art of technology, and how those effects may ripple out to change society at large. While these analyses would be much simpler were it not the case, the fact of the matter is that technological innovation is part and parcel of human society. Advances in technological innovation cannot be extricated from the social conditions which gave rise to such innovation, nor can they be separated from the social consequences which result from the adoption of that technology. If we wish to base our analyses in historical fact, it seems prudent to first review some basic history with an eye towards the factors we intend to consider. If we wish to see a preview of the impacts that automation will have on the workforce, and consequently the society at large, it seems prudent to look at the ways automation shaped society as its effects on the economy were first made manifest.

Let us consider, for a moment, the steam hammer. The steam hammer, a powered hammer for forging and stamping of metal parts, was invented in the mid-1800s as increasingly large and complex machines necessitated the fabrication of larger and heavier components. It made routine the construction of many works that, only years before, would have been considered marvels of engineering and concerted human effort. It contributed to the construction of great ships, the production of delicate clockwork, and a massive increase in the industrial productivity of Western Europe that saw an increase in the availability of cheaply manufactured goods formerly restricted to only the richest members of society. It also contributed [14] to a sharp decrease in the leverage that skilled engineers held over their employers, resulting in drastic consequences for working environments, and sharp increases in child-labor. James Nasmyth, one of those credited with the invention of the steam hammer, made the following statement regarding the power of automation in the modern (as of 1851) factory:

The characteristic feature of our modern mechanical improvements is the introduction of self-acting machinery. What every mechanical workman has now to do, and what every boy can do, is not to work himself, but to superin-

tend the beautiful labour of the machine. The whole class of workmen that depend exclusively upon their skill is now done away with. Formerly, I employed four boys to every machine. Thanks to these new mechanical combinations, I have reduced the number of grown-up men from 1,500 to 759. The result was a considerable increase in my profits. [15]

From the specific, we proceed to the general. The example of the steam hammer (in particular, the contrast between its positive and negative social effects) raises one of the first difficulties we will be forced to surmount in the course of developing our ethical framework. In the case of the industrialization of the workforce of Western Europe, those who wished to deploy greater degrees of automation in their factories were directly at odds with those who work in those factories. This is reflective of the effects of automation, not just on the market for the product they create, but also on the labor market. Labor, like any other commodity, has a price, which fluctuates based on supply and demand. When the owners of the factory can produce just as much value from less work, the demand for that labor goes down, and its price drops. Those workers have less leverage, and, as observed by Gaskell, their working conditions suffer.

When balancing the gains in technology and standards of living that have resulted from this march of progress, we see that there are two sides of this coin we must consider. The society as a whole may benefit, but what of the workers that toil in the factory? Does the gain outweigh the harm? Fortunately, in this particular case, we are not forced to make such a decision: social and cultural pressures intervened. The labor movement, in an internationally coordinated effort, slowly but surely pushed many governments into enacting legislation that protected workers from unfair practices on the part of their employers.

And yet, we are faced with a somewhat tougher questions, when we step back and consider the broader context around the industrialization of Western Europe. Industrialization is a somewhat titanic effort, it requires the investment of massive amount of capital to begin to see the advantages resulting from leveraging the economies of scale that make automation profitable. Different societies have achieved such a concentration of wealth in different ways. China, in the course of The Great Leap Forward, began industrialization on a massive scale by nationalizing all production, taking central command of labor and distributing it according to a grand plan for industrialization of the workforce as a whole. The Soviet Union took a similar approach, nationalizing industry and placing it under central command, but then decentralized the industry geographically, exporting raw materials to nearby states as it absorbed them, thus creating a demand for the growth of industry in those formerly independent states that drove their growth further outward until they started bumping up against NATO countries. Western European countries accumulated wealth in privately held trading companies by establishing an early foothold in mercantile capitalism through the Triangle Trade. That is, they establishing trade routes between Europe, Africa, and the Americas, alternating in the shipment of slaves, raw materials (such as sugar cane grown by slaves), and manufactured goods (such as rum produced from last year's sugar cane).

In each of these situations, there exists a contradiction between the forces acquiring enough wealth to assemble an industrial base, and the people that had to give up that wealth. In China and the Soviet Union, the costs were paid by the citizens of those countries, and they reaped the benefits slowly as their standards of living increased fairly evenly. In Europe, the standards of living and availability of luxury goods skyrocketed much more rapidly: after all, the cost of automaton was offloaded to somewhere else: Africa and the American colonies. In the Soviet Union and China, the analysis looks simple: did the country gain more than it lost? China suffered a number of famines as the Great Leap Forward encountered organizational difficulties, but over a period of 30 years, the average lifespan of a Chinese citizen nearly doubled.

When, however, we look at a situation like British industrialization, we now find that we must take a look at a much more diverse selection of groups in analyzing the impact of technology on the society. As globalization has blurred the lines between nations, so too has information technology

led to a vastly more interconnected world than we have ever seen. A change in technology today can lead to nearly instantaneous changes through enormous swaths of society. In the immediately following sections, we will further examine the current state of the world as it has evolved, the approaches that some groups take in analyzing that world, and the steps they deem necessary due to such an analysis.

III. THE CURRENT SCENARIO

On the 12^{th} of March 2016, history was made. Lee Sedol, a professional Go Player who some consider to be the best player of the last decade, lost a game of Go. His third consecutive loss in a best-of-five competition. His opponent? AlphaGo an Artificial Intelligence based Go playing computer system [16]. Although this wouldn't be considered anything new, with computers outsmarting their human opponents from Watson, IBM's Artificial Intelligence besting previous champions to win Jeopardy [17] to Deep Blue(another of IBM's attempts at Artificial Intelligence), the first computer to beat a Grand Master at Chess. But why was this such a historic victory? It was for two primary reasons; firstly the game of Go has about the same number of possible positions as there are atoms in the universe (this rules out simply computing every possible combination) and secondly, this is a milestone in the trend of using computers to replicate mental processes that were earlier considered to be the sole dominion of the living beings, particularly humans.

As mentioned earlier, we live in the Information Age, and changes have a tendency to propagate through vast sections of society fairly quickly. Globalization and the subsequent transfers of technology enabled by the Internet, allows for a proliferation of ideas and techniques that would have been unheard of a few decades ago. Automation's strongest argument has been the one to reduce manual labor and improve the quality of human lives. Although now, it would seem that this is no longer the case. Many proponents of automation have pointed to the migration of people from unskilled and semi-skilled labor to more skilled labor which

rely on mental processes that could not have been automated at the time. It is evident that this sacred boundary is under threat of being or has already been, depending on who you ask, infringed upon.

The advent of Applications that have improved accessibility to services, from planning your week to simply setting an alarm or a reminder to do your taxes, all these tasks have been automated already. The automation of these activities begs the question whether this would make humans obsolete in performing these activities. The public sentiment does not seem to reflect this. In a study carried out by the *Pew Research Center* [18], although about two-thirds of Americans believe that computers and robots will be performing most of the work being done by humans in the next 50 years, nearly 80 % believe that their current jobs will still be around then.

A likely answer to this is skewed figure is probably to do with the General outlook most people have on technology. In a 1996 *Washington Post/Kaiser Family Foundation/Harvard* poll, 70% of Americans said the increased use of technology in the workplace was good for the economy. In a 2010 *Allstate/National Journal* poll, 79% said that information technology was extremely or very important to creating economic growth in the U.S. Some would say they aren't wrong but it should be noted that it is rare that general public opinion is aligned with the most ethical option.

Even though AI forms only a small section of this problem, it engenders the primary difference between the *First Industrial Revolution* and now, the fact that mental processes can be performed far more efficiently by computers. Technology Heavyweights like Elon Musk and Stephen Hawking have openly spoken about the perils of the technology when it is used for warfare [19]. With some expressing concern about the possibility of it being the cause of human extinction. The extreme views on this issue are well documented, and we will further elaborate on them in the next section, but our aim remains to balance these views to ensure a smooth transition. There is no way to , the best we can do is ensure it is done in the most ethical fashion, without favoring one stakeholder over

the other and ensuring that we have a holistic view of the situation rather than proposing solutions that would be suited to a vacuum.

IV. A Review of Present Solutions

A. The Hope of the Singularity

B. Bill Joy

C. Jason Lanier

D. The Amish

The Amish in most popular depictions are painted as rural farmers. They are often protrayed as a culture who fear technology. However, this is a gross simplification of the relationship the Amish have with technology. Previously we have seen Bill Joy who is worried about the juxtaposition of human nature, underestimating the probability of catestrophic effects, and technology and Jason Lanier who critiques what he sees as the rampant corporatism with technology. In contrast, the Amish are not so worried about negative effects like pollution, but the effect of new technologies on their culture. The Amish view of technology is that technology and society or culture are created together and in the words of Sheila Jasanoff "both embeds and is embedded in social practices,identities, norms, conventions, discourses, instruments, and institutions " [20]. The Amish regard technology as a force that has the power to reinforce social norms, control how people interact, or enable new forms of interaction between people. Where Bill Joy implore us to accurately weigh the net harm including the likelihood of catestrophic events with respect to technology and Jason Lanier insists that we need to insure that corporations do not gain most of the benifit of new technologies instead the Amish insist that any new technology is in line with their culture as well as reinforces their tradition, religion, and culture regardless of any objective potential benefits or harm associated with that technology.

E. Deep Green Resistance

Deep Green Resistance is a group primarily identified by their beliefs and behavior with regard to specifically environmental matters, rather than more general opinions about technology. Specifically, DGR subscribes to the notion that industrial civilization poses too great of a threat to life on the planet earth, both human and otherwise, to be allowed to survive. The ultimate goal of DGR is the destruction of industrial civilization and a return to an earlier stage of societal and economic development.

DGR's philosophy draws from the Deep Ecology movement, which holds that an anthropocentric analysis of ecology, defining other forms of life in terms of their utility to humankind, does not do an adequate job of describing the complexity of ecological systems. DGR takes this a step further, in declaring all life to be equal to human life. Starting at the problem of inequality among life, they turn to ideals similar to the political stance of anarcho-primitivism: the idea that human society, industrial capitalism in particular, cannot be reformed into a more beneficial form.

Such a philosophy stems from an agreement with Friedrech Engal's seminal anthropological work on early civilizations [21], specifically the assertion that early gatherer-hunter societies were unable to produce structural inequality without the ability to acquire surplus wealth. Unlike Engals, modern Marxists, and even their fellow anarchists, DGR and anarcho-primitivists do not believe that a society can ever reach a level of egalitarianism and freedom from oppression that they are comfortable with. Rather, they believes that the only truly equal societies that have ever existed were only able to function without the trappings of civilization. Therefore, they are willing to advocate any means to right the wrongs they see with the world. To them, technology is not evil per se. Rather, technology is another product of a system with which they cannot make peace, and as such, it must fall by the wayside.

Later in our paper, we advance an ethical framework for the analysis of decisions regarding technology and automation. In so doing, we emphasize that we must look to the past when examining possible consequences of an action, basing, wherever we can, our reasoning on historical fact. In this case, we would be remiss if we did not qual-

ify our description of the views of DGR with the caveat that their views ultimately reduce to something chillingly familiar. A return to pre-industrial levels of technology would necessitate either the death of billions, or the restriction of new births to a fraction of their current levels. Without modern medical technology, it becomes a disturbing fact that disabled people will not be able to survive in such a world: this means DGR endorses either genocide or eugenics.

V. A Coherent Ethical Framework

A. *Identifying Stakeholders*

Ethical Analysis of an action, circumstance or rule must always begin with Stakeholder Analysis. Stakeholders are those who are affected as consequence of the action, circumstance or rule. In other words, the parties concerned have a stake in the eventual outcome of performing the action or implementing the rule.

The framework aims to take into account two main factors; firstly, factors that affect the identification of stakeholders; and secondly, when viewing automation we observe three general classes of stakeholders, with varying levels of interplay seen between these classes based on the aforementioned factors.

Factors affecting Stakeholder Identification:

The process of identifying Stakeholders is generally specific to the action and rule as well as to the circumstances that surround the process. Although our framework strives to espouse the Kantian Ideal of universality, it would be naive to assume that the stakeholders would be a constant across the board; taking into account that our framework aims to perform an ethical analysis on the effect of automation To this end, in order to identify the stakeholders we must try to obtain a holistic view of the situation. In order to do this, we enumerate three factors, which would have an effect in the identification process:

- Cultural Bias:

    Cultural Bias has long hindered the adoption of technology. Rather than attributing this as a drawback (which it is sometimes) see it as nothing but the decision making process of

a collective, based on previous experience as well as skepticism about new and sometimes untested methods/technologies. Like a double-edged sword, the same cultural bias might catalyze the adoption of the technology as well. This can be seen in the early problems faced during adoption of disruptor technologies like the Radio and aircrafts. In order to correctly identify stakeholders, it is imperative to accurately weigh the cultural bias associated with a given society; especially since the effects of the technology might be experienced only by a certain section of society.

- Value System:

    Value systems play a big part in helping the acceptance of technology. Say for example, in the United States the Apple vs. FBI debate has tested the value systems of the public, where protection of individual liberties is of paramount importance, if a similar case would have taken place in the People's Republic of China, the prevailing value systems would have ensured massive public support for the Government's stand. In short, the classification of stakeholders would change dramatically where in one society, the Government represents the interests of the public where as in another, the individual/public takes on an identity of their own.

- Economic Status:

    The Economic Status of a society broadly refers to the distribution of wealth in the society. We see vastly contrasting structures in Developed Countriies and in Developing Countries. One example that could elucidate this difference is the adoption of automation of jobs such as customer service. This would have varying effects depending on the country. In developed countries, these are considered tertiary level jobs and do not form the backbone of the economy. Hence the effect would be largely mitigated and with government intervention could also be avoided. But in the case of developing countries, where semi-skilled jobs like freight services employ a large number of people, we would see a larger effect in

terms of unemployment. Accounting for this would play a large part in correctly weighing each stakeholder.

B. General Classes of Stakeholders

Taking into account the effect of automation over the entire cross-section of society, we think it is prudent to identify certain classes of stakeholders which would be present across any situation, regardless of the geographical location, economic status and value system. Although there would be different interplays between these classes given the above mentioned factors that affect the selection of stakeholders, their existence is no less in doubt by our standard.

It is important to add that although we define these classes, they are in no way exhaustive and the due diligence required to identify stakeholders must be done in order to ensure the holistic view we would like to achieve in this part of the process.

- Public:

    The Public, refers to society at large who no matter what the situation will be directly/indirectly affected by the implementation of the automation. Typical issues include job/labor concerns, concerns regarding misuse against other members of the public, and concerns regarding laws enforcing the ethical use of the same, just to name a few. By far this is one of the larger classes, and within it many more might emerge, especially when the technology is targeted to a certain section/strata of society. It is therefore essential to identify which section might have a larger stake in the implementation of the automation, be it positive or as it is often seen, negative.

- Government:

    Government, refers to the regulatory body that oversees the ethical use of the automation. We can say in certain scenarios, the lines discerning public interest and government interest might blur, we considered these to be exceptional cases where we can identify one class to represent the interests and stakes of another(as mentioned in the example above). It can be seen in most modern societies today, the public does not necessarily gain adequate representation of their interests through government and similarly, the Government must at times ensure the security and well-being of the state even if that does not represent what the public deems fit. For this and reasons that outline similar cases of dichotomy, we would like to represent this as a separate class.

- Corporate Interests:

    Automation today, as it was earlier as well, is driven by commercial interests. We represent these interests through the class of corporate stakeholders. These would represent the individuals as well as organizations which have driven the case for the automation and would benefit in some monetary fashion by the said implementation. Apart from this, we would like classify industries in this class of stakeholders as well, because of the widespread transmission of ideas and technology, practices linked with an increase in productivity would be widely adopted even if that is at the cost of jobs. Although it would be easy to classify corporate interests as strictly monetary, but in certain scenarios we see that this stakeholder has aligned interests with another classes(such as Government) and hence would require a somewhat different approach to weighing their stakes.

C. Defining a Coherent Value System

Subsequent to identifying stakeholders and classifying the stakeholders into the basic classes, or new ones depending on the circumstances, we must elucidate upon a value system that would act as a guide and in the process allow us to reach a conclusion regarding the ethical implementation of the automation in question.

We propose a negative utilitarian framework focused on preventing harm, where we define harm in terms of violation of a number of values that we consider to be important. We apply this framework with the caveat that its usefulness may break down at certain moral extrema, which result in a reluctance to make trade-offs between particular harms in particular consequences. We identify a rough

taxonomy of harms as violations of the following values for stakeholders in an ethical analysis:

1. Autonomy of Self

    Autonomy of self, refers to the basic human rights as laid down by the Universal Declaration Human Rights [**?**], which includes but is not limited to, Right to life, Right to freedom of expression

    As laid-down in the Charter, any harm to these basic rights can considered a violation of our framework, and the person responsible for the analysis should eschew any automation that endangers these basic rights, or in the worst case to mitigate these violations through implementation of other practices that would ensure protection of these rights to the highest degree possible.

    A vital point within this tenet is the Right to Privacy, although most people believe this is a fundamental right, in a majority of cases this would take precedence over any technology that infringes on this right. As we are seeing the battle rage between Apple and the FBI, the protection of this entire section of rights would come under the social contract between the individual/s and society at large.

2. Welfare/Well-Being

    The ability to ensure Welfare and well being of self, family and society at large has long been considered the highest responsibility. The world we live in today, this is directly associated with sustenance as well as health.

    Any violation of the welfare of the stakeholders through either contamination of sources of sustenance or loss of livelihood could be considered a violation of this principle. Understandably, there are costs to progress, although this framework aims to minimize the infringements, we propose that on finding a violation, action should be taken against the party to ensure that the cost of implementing the automation would be higher than the gain from such technologies, which are implemented by violating these rules.

    The purpose of this framework is less to take a moral stand but rather a way to balance the effect of consequences on the stakeholders, by suggesting action we believe that the framework used in conjunction with the ground reality would help the user come to an understanding of how to ensure a balance is reached.

D. Synthesis

VI. Self-Driving Cars: an Application of our Framework

A. Background

Long a fascination for inventors, the seed for autonomous cars was sewn in the 1960s with the Stanford Cart [22]. The cart navigated small spaces by taking photos of its surroundings, using a simple computer program to analyze the photos, and choosing a path to travel depending on that analysis. Two decades later, a vision-guided Mercedes-Benz van designed by Ernst Dickmanns navigated empty streets at speeds of up to 39mph. By 2009, Google began privately testing anonymous vehicles on closed streets, and by 2012, changes in various states' laws allowed for testing on city streets. As of 2016, Google claims over 1.3 million autonomous miles driven, although Google's cars are limited to 25mph and must always have a human driver present. [23] As driverless car technology improves and becomes more widely implemented, the true societal impact of this technology will begin to be felt.

B. Application

We're in an interesting time regarding autonomous vehicle technology. It's clear that the technology will become widely used, but it has yet to actually happen. This presents an opportunity to consider how, as a society, we should approach the ethical ramifications of this potentially society-altering technology. To do so, we'll apply the ethical framework laid forth earlier in this paper to autonomous cars.

To begin, it is important to identify the three main stakeholders in driverless car technology: the public, the government, and corporations. These stakeholders are universal, but the beliefs of each stakeholder can vary depending on cultural biases,

value systems, and economic status. Also, the weight given to each stakeholder can vary based on these factors. In the interest of concision, this application will focus on a generalized idea of American cultural bias, values, and economic status.

Because the United States is such a demographically, economically, and culturally diverse country, it is difficult to make generalizations of the country as a whole. That being said, it is possible to extrapolate mainstream value systems of our three stakeholders (public, government, and corporations). The American public values individualism and personal rights. Corporations value free market principles and the acquisition of capital. The government acts as a mediator (sometimes poorly) between the interests of corporations and of individuals. Each stakeholder has the power to act on the other and bring about change. With the stakeholders, value systems, and power structure identified, it is possible to apply our ethical framework to autonomous vehicle technology in the United States.

Autonomous vehicles have a lot to offer a utilitarian. A recent Virginia Tech study [24] shows that the crash rate of Google's self-driving cars is lower than the national crash rate of conventional cars. Google's cars experience 3.2 crashes per million miles while conventional cars experience 4.2 crashes per million miles. With autonomous vehicle technology continually improving, the technology will become even safer in the years to come. This sounds great, but is a full on, no-strings-attached embrace of the technology the best course of action? For a utilitarian, the answer may be "yes", but there is a problem with this approach. Autonomous vehicles will face a variety of moral dilemmas, and a one-size-fits all ethical framework doesn't cover all of these possibilities sufficiently.

Consider the following example: during the November 2015 terror attacks in Paris, ride-sharing app Uber's algorithmic design dramatically raised prices in the areas near the attacks due to an increase in ride requests from people trying to flee the chaos. In this situation, stakeholders (public and corporate) were in direct opposition to each other. Uber saw the situation as a moneymaking oppor-

tunity, while the public saw Uber as price gouging and profiting from a disaster. Eventually, Uber gave in to public demand and lowered prices. Situations such as this will become the norm in the age of autonomous cars. In fact, Uber is aggressively seeking driverless car technology for use in their taxi company.

In the case of Uber raising the prices on people fleeing a disaster, utilitarianism supports profit maximization as a means of supporting the greatest societal good. Clearly, this is a morally unfavorable position for the public, but fortunately, the moral framework set forth in this paper allows an exception to utilitarianism on the basis of protecting an individual's welfare and well-being. In the case of the Paris attacks or hypothetical future situations involving Uber autonomous vehicles, our moral framework prioritizes lives in danger over profits in dangers.

Consider another hypothetical example: you're riding down the highway in your autonomous car and a large object falls off of a truck in front of you. Your car cannot stop itself in time and must decide on the best course of action. If it continues straight and runs into the object, it places your life in serious jeopardy. There are two motorcyclists on either side of your car; one is wearing a helmet and the other is not. Swerving into one of the motorcyclists will likely save your life, but jeopardize the life of one of the motorcyclists.

With strict utilitarianism, your car would likely swerve into the motorcyclist wearing the helmet because it creates the best chance of survival for all involved, but is it fair to punish a motorcyclist for being responsible and wearing a helmet? What are the legal and moral ramifications of such a programmed decision? In this scenario, our ethical framework would fall back onto the autonomy of the individual. They could decide how their car is programmed to react in cases such as these. Alternatively, self-driving cars could rely on an element on randomness in making such decisions or even cede control to a human driver (local or remote) despite the car's ability to handle situations such as this without human intervention. The important point is leaving autonomy in the hands of the indi-

vidual, not the machine.

While far from foolproof, our ethical framework allows for important exceptions to utilitarianism. The implementation of driverless cars will undoubtedly lead to many unique moral dilemmas, therefore the existence of *fallback rules* to a society's moral framework will be critical. In the previous examples, the *fallback rules* of our ethical framework placed health and well-being above financial profit and the autonomy of a human over the autonomy of a machine. These are only two examples of the multitudes of moral dilemmas that will become ever-present in a driverless car future, and it is critical that we continue to work toward developing an ethical framework to deal with every possibility.

VII. CONCLUSION

Rise of the automated tools designed to accomplish certain tasks and ethical problems stemming thereof is a very potent area for dialogue. It is evident that the set of current solutions have more cons than pros when one put them under scrutiny. We claim that the framework that we proposed in this paper would minimize the short comings and maximize the gains. But one must always remember that until we do obtain a certain level of automation, real the social impact of such a technology can only be speculated. Minute details that escape the most prudent analysis might have very far reaching consequences. And thus we stand on the precipice and stare in to the void, and does the void stare in to us back?

REFERENCES

[1] Homer, *Iliad*, 1194 BC, Chapter 18. 371 ff.
[2] Aaron J. Atsma, "Automatones," 2015.
[3] Lie Yukou, *Liezi*, pp. 83–86, 600 BC.
[4] Sarah L. Higley, "Alien intellect and the roboticization of the scientist," *Camera Obscura*, vol. 14, no. 1-2 40-41, pp. 129–160, 7 1997.
[5] "Google self-driving car project," 2016.
[6] "Your autopilot has arrived," 2016.
[7] I. Liu and B. Ramakrishnan, "Bach in 2014: Music Composition with Recurrent Neural Network," *ArXiv e-prints*, Dec. 2014.
[8] James Gould-Bourn, "Man combines random peopleŠs photos using neural networks and the results are amazing," 2016.
[9] Justin Mahboubian-Jones, "Fully-automated security robots have arrived," 2016.
[10] Eyal Reingold, "The Turing Test," 1999.
[11] Arthur C. Clarke, *Profiles of the Future*, chapter "Hazards of Prophecy: The Failure of Imagination", Phoenix, 1962.
[12] Michael Sainato, "Stephen hawking, elon musk, and bill gates warn about artificial intelligence," 2015.
[13] Luke Muehlhauser, "What is agi?," 2013.
[14] Peter Gaskell, *The manufacturing population of England, its moral, social and physical conditions*, Baldwin & Cradock, 1833.
[15] Karl Marx, "The strife between workman and machine," in *Capital*, vol. 1, chapter 15. 1867.
[16] Cade Metz, "GoogleâĂŹs ai takes historic match against go champ with third straight win," 2016.
[17] John Markoff, "Computer wins on âĂŸjeopardy!âĂŹ: Trivial, itâĂŹs not," 2011.
[18] Aaron Smith, "Public predictions for the future of workforce automation," 2016.
[19] Samuel Gibbs, "Musk, wozniak and hawking urge ban on warfare ai and autonomous weapons," 2015.
[20] Jameson M Wetmore, "Amish technology: Reinforcing values and building community," *Technology and Society Magazine, IEEE*, vol. 26, no. 2, pp. 10–21, 2007.
[21] Friedrich Engels and Tristram Hunt, *The Origin of the Family, Private Property and the State*, Penguin UK, 2010.
[22] Hans P Moravec, *The Stanford cart and the CMU rover*, Springer, 1990.
[23] Alex Davies, "GoogleâĂŹs self-driving car caused its first crash," *Wired*, 2016.
[24] Myra Blanco, Jon Atwood, Sheldon Russell, Tammy Trimble, Julie McClafferty, and Miguel Perez, "Automated vehicle crash rate comparison using naturalistic data," Tech. Rep., Virginia Tech Transportation Institute, 2016.